

Next-Day Fatigue Prediction Using Apple Watch Physiological Data: An Edge Computing Approach

Pratik Patil
George Mason University
CS 655
G01523747

Ritesh Sanjay Sandbhor
George Mason University
CS 655
G01525006

ABSTRACT

Apple Watch and HealthKit collect sleep, heart, activity, and respiratory signals that can be useful for fatigue prediction. A common issue is that these signals are personal, so sending them to a cloud service is not ideal for a privacy-focused mobile application. In this project, we built an iOS prototype that predicts next-day fatigue directly on the phone from Apple Watch-derived HealthKit data. The offline pipeline parses Apple Health exports, creates night-level sleep and activity features, trains classical machine learning models with chronological splits, and exports a compact JSON inference contract that Swift can execute. We evaluated the pipeline on 359 labeled nights. The best reported baseline, Random Forest, reached 52% accuracy and 0.51 F1 on the hold-out test set. This result is not strong enough for production use, but it shows the difficulty of predicting subjective fatigue from wearable data alone. The main outcome is a working mobile path: the app requests HealthKit permission, builds the feature payload, fills missing values, and computes the prediction locally without a cloud service.

KEYWORDS

Fatigue Prediction, Apple Watch, HealthKit, Edge Computing, Machine Learning, Wearables, Data Privacy

1 INTRODUCTION

Apple Watch records many signals that are related to recovery, including sleep stages, heart rate, HRV, SpO₂, steps, and active energy. These measurements are useful, but they are also sensitive health data. For a mobile computing project, this makes fatigue prediction a good case study: the model should be useful, but the data should stay on the user's device.

Fatigue affects attention, productivity, and daily decision making. Clinical measurement can require surveys or controlled testing, while wearable systems estimate recovery from motion and physiological signals. Many machine learning systems still depend on centralized processing, which adds privacy and data-handling concerns for biometric data [1].

This project asks whether a class-scale fatigue prediction system can run locally on an iPhone. The system has two

parts: an offline Python training pipeline and a Swift HealthKit application. The Python side prepares features and trains the models. The Swift side loads a small JSON contract, fetches HealthKit data, prepares the same feature order, and runs the prediction without a server request.

The report focuses on four implementation pieces:

- (1) Parsing Apple Health exports into night-level records with sleep, HRV, heart-rate, oxygen, step, and activity features.
- (2) Evaluating models with chronological train/test splits so later days are not used to predict earlier days.
- (3) Exporting the trained mobile model into a JSON contract that includes feature order, imputation values, scaling values, coefficients, and threshold.
- (4) Building a Swift HealthKit demo that can request permission, load recent data, and compute the fatigue prediction with no network dependency.

2 BACKGROUND AND RELATED WORK

2.1 Actigraphy and Heart Rate Variability

Sleep analytics often start with actigraphy, which uses movement patterns to estimate sleep and wake periods [2]. Apple Watch also provides heart-rate and HRV-related measurements from photoplethysmography (PPG). HRV is often used as a rough indicator of autonomic nervous system recovery [5]. In this project, we use Apple Watch sleep stages, heart rate, HRV, respiratory rate, and SpO₂ as the main nightly signals. These signals are noisy, so the pipeline uses summary statistics instead of assuming each raw sample is reliable.

2.2 Mobile Edge Computing Paradigms

Edge computing moves computation closer to where data is created. For wearable data, this means the phone can process the user's signals instead of uploading them to a remote server. Federated learning and Edge AI are related approaches that try to reduce centralized data collection [3, 4].

For this project, mobile constraints were treated as design requirements rather than afterthoughts. The deployed inference path therefore avoids server calls and runs as direct numeric computation in Swift. This keeps the prototype small

enough for an iPhone demo while preserving the local privacy goal.

3 SYSTEM ARCHITECTURE

The system has three parts. The offline Python pipeline prepares the dataset and trains models. The VitalSync web dashboard helps visualize trends during development. The iOS application is the main mobile demo and runs the final prediction path.

3.1 Data Engineering and Pipeline Architecture

The offline backend is a Python pipeline that works from Apple Health export files. Apple Health exports are XML files, so the first step is to parse `export.xml` and extract the relevant `<Record>` entries. For this project, the important record types are Heart Rate, HRV, Respiratory Rate, SpO₂, Sleep Analysis, steps, and active energy.

3.2 VitalSync Frontend Dashboard Design

The VitalSync web dashboard was used as a development and presentation aid. It shows sleep trends, HRV changes, and predicted fatigue labels in a browser. This helped us inspect the data and explain the project before relying only on the physical iPhone demo.

3.3 iOS Native Edge Deployment

The iOS app is written in Swift and uses HealthKit for local data access. Instead of shipping Python code or a large ML runtime, the app loads `mobile_linear_contract.json`. This file stores the feature order and model parameters needed for inference. The app then fetches recent HealthKit samples, computes the required features, applies the same pre-processing used during training, and returns the fatigue prediction on the phone.

4 SYSTEM IMPLEMENTATION AND DATA PROCESSING

4.1 Wearable Sensing and Night Windows

The first data-processing problem was assigning sleep to the correct day. Sleep usually starts before midnight and ends after midnight, so grouping only by calendar date would split one sleep session into two days. We used a “Night Window” rule with a 6-hour time shift. With this rule, a night that starts on Day T and ends on Day $T + 1$ is stored under the Day $T + 1$ label.

4.2 On-Device Data Transformation

After creating the night windows, each night becomes one row in the training table. We use summary statistics because

Table 1: Predictive Performance on 30% Test Set

Model	Acc.	Prec.	F1
Logistic Regression	49.0%	0.48	0.47
Random Forest	52.0%	0.50	0.51
Gradient Boosted Trees	50.9%	0.49	0.50
Support Vector Machine	48.1%	0.47	0.46

wrist sensors can be noisy and irregular. The main feature groups are:

- (1) Total Sleep Duration (Float Hours)
- (2) Median HRV (ms) and Standard Deviation HRV
- (3) Minimum, Maximum, and Median Resting Heart Rate (BPM)
- (4) SpO₂ minimum limits, reflecting brief anoxia drops.
- (5) 7-day rolling step count mean and standard deviation
- (6) 7-day rolling active energy expenditure mean

We also added rolling multi-day metrics. The reason is that fatigue may build over several days, so a 3-day or 7-day trend can be more useful than one isolated night.

4.3 Ground Truth Collection

The final dataset contains 359 labeled nights. The labels are binary: 0 means focused, alert, or normal, while 1 means high subjective fatigue. After merging labels with the feature rows, the dataset contains 180 non-fatigued nights and 179 fatigued nights.

5 SYSTEM EVALUATION IN REALISTIC SETTINGS

5.1 Evaluation Design and Predictive Capability

Random K-Fold validation is not ideal for this dataset because it can mix future days into the training set and earlier days into the test set. That would make the evaluation look better than a real deployment. Instead, we split the dataset chronologically into 70% training data and 30% hold-out test data.

We evaluated several classical models that are small enough to inspect and compare. Random Forest gave the best reported result in this run, with 52.0% accuracy and 0.51 F1-score. The result is only slightly better than chance, so it should be treated as a feasibility result rather than a reliable fatigue detector.

5.2 On-Device Inference Contract and Runtime Behavior

For mobile inference, the final model path is exported as `mobile_linear_contract.json`. At prediction time, the iPhone

app builds a feature dictionary from HealthKit-derived nightly signals. If a feature is missing, the app uses the training median stored in the contract. It then scales the values and computes the class score with vector arithmetic.

This format is simple, but that is the point for a mobile demo. The app does not need to initialize TensorFlow Lite, Core ML, or a Python runtime. In our test, the path from HealthKit data fetch through final prediction display completed in under 200 milliseconds on an iPhone 14.

5.3 Confusion Matrix and Feature Importance Analysis

The Confusion Matrix (Figure 1) shows that the Random Forest model correctly identified 39 out of 55 actual fatigue events. However, it also produced 37 false positives. This means the model often predicts fatigue when the label is non-fatigued, which is one reason the current model should not be presented as production-ready.




Figure 1: Confusion Matrix illustrating model classification on the hold-out set.

The Feature Importance plot for the deployed mobile contract (Figure 2) was regenerated from the same inference contract used by the iOS app. Importance is reported as the magnitude of each feature’s coefficient on standardized inputs, i.e. the per- 1σ change in log-odds. Blue bars push the prediction toward the fatigued class, while red bars push it toward the alert class.

The strongest features are mostly rolling sleep metrics. The 7-day rolling deep-sleep mean is the strongest stabilizing feature, while 7-day rolling total sleep mean and 7-day

rolling REM mean are also important. Heart-rate variation and respiratory-rate mean also appear near the top. This supports the idea that the model is using accumulated sleep and recovery context rather than only one night’s data. One issue we noticed is that *sleep efficiency* contributes almost no importance because it was nearly constant in this dataset.




Figure 2: Feature importance for the deployed mobile inference contract: $|\text{coefficient}|$ on standardized features (per- 1σ log-odds). Blue = positive (toward fatigued); red = negative (toward alert).

6 SYSTEM DEMONSTRATION OF EDGE COMPUTING CAPABILITIES

Figure 3 shows the iOS screens used to validate the system path. The screenshots cover the app interface, HealthKit-driven data loading, local prediction output, and the history view used during the demo.

During the demo, the user grants HealthKit access, loads Apple Watch-derived data, runs local inference from the bundled contract, and reviews the fatigue result. This confirms




Figure 3: Screenshots from the system demo showing the app interface, local prediction output, and on-device history view.

that the project connects the offline model pipeline to a native mobile workflow instead of remaining only an offline notebook result. All inference is performed locally, so the demo does not require a network connection or a cloud prediction service.

7 DISCUSSION AND LIMITATIONS

7.1 Mobile Runtime Constraints

The iOS prototype also showed practical runtime limits. HealthKit access requires user permission, real data is only available on a physical iPhone, and refresh behavior must respect iOS scheduling rules. For the class demo, the app uses manual fetch and prediction as the primary path, which makes the privacy behavior and data flow easier to verify.

7.2 Addressing the Performance Ceiling

The model performance is limited by the small and mostly single-user labeled dataset. A 52% accuracy result should therefore be interpreted as a feasibility baseline, not as a production-quality fatigue classifier. Fatigue is subjective and can be affected by stress, workload, illness, and schedule changes that are not fully visible in sleep and vital-sign data. The main value of the project is the working mobile pipeline: data access, feature extraction, local inference, and UI feedback all run in one privacy-preserving path.

7.3 Implementation Lessons

The implementation clarified that a wearable ML model is only useful if the preprocessing path is reproducible on the device. The same feature names, ordering, imputation policy, and scaling parameters used during training must be available to Swift at runtime. Otherwise, the phone may compute values that look valid but no longer match the training distribution.

A second lesson is that HealthKit data is not a clean table. Samples arrive from multiple sources, some records are user-entered, and simulator data cannot replace a physical iPhone for real validation. The prototype handles this by preferring Apple Watch-origin samples when present, filtering manual entries where possible, and making missing-feature behavior explicit through the contract.

7.4 Fixed Sleep Windows vs. Dynamic Rhythms

A limitation of the current pipeline is the fixed night-window rule. The 6-hour shift works for a normal overnight sleep schedule, but it may fail for shift workers, naps, or irregular sleep. A better version would calculate sleep windows from the actual sleep samples instead of relying on a fixed time offset.

7.5 Future Directions: Rolling Context

The next step is to improve the rolling features. In particular, 3-day and 7-day HRV trends may capture recovery better than one-night summaries. More labeled users would also be needed before the model could be evaluated as a general fatigue predictor.

8 CONCLUSION

This project built and tested a mobile-first pipeline for next-day fatigue prediction from Apple Watch data. The work covered data parsing, night-window construction, feature engineering, model evaluation, JSON contract export, and Swift HealthKit integration. The model result is modest, with 52% accuracy and 0.51 F1, so the current system should be viewed as a prototype rather than a finished fatigue classifier.

The main learning from the project is that mobile ML is not only about choosing a model. The training pipeline,

feature contract, missing-value handling, HealthKit permissions, and iOS runtime constraints all have to match. Future work should focus on more labeled users, better sleep-window detection, stronger rolling recovery features, and a longer real-device evaluation.

REFERENCES

- [1] Q. Chen, Y. Wang, and X. Liu. 2021. Privacy-Preserving Edge Machine Learning: A survey. *IEEE Communications Surveys & Tutorials* 23, 4 (2021), 2345–2366.
- [2] A. Martinez and J. Smith. 2022. Actigraphy and smartwatches in commercial sleep tracking. In *Proceedings of the 2022 ACM International Joint Conference on Pervasive and Ubiquitous Computing*. ACM, 112–120.
- [3] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas. 2017. Communication-efficient learning of deep networks from decentralized data. *Artificial intelligence and statistics* (2017), 1273–1282.
- [4] R. Raskar and V. Chopra. 2020. Edge AI: Deploying advanced machine learning models safely to constrained devices. In *ACM Symposium on Edge Computing*. ACM, 34–42.
- [5] F. Shaffer and J. P. Ginsberg. 2017. An overview of heart rate variability metrics and norms. *Frontiers in public health* 5 (2017), 258.